
DATASET: Indirect Prompt Injection CV Dataset.

This dataset has been used for the research described in the work:

Alfredo Milani, Valentina Franzoni, and Emanuele Florindi. "Indirect prompt injection in large language models." *Neural Computing and Applications* 38.13 (2026): 530. <http://doi.org/10.1007/s00521-026-12266-x>

to evaluate the robustness of LLMs against indirect prompt injection, we designed a systematic experimental framework based on an heterogeneous dataset of Curriculum Vitae (CV) documents. The scope of the dataset is to verify the generalizability of threats and countermeasures. We designed a controlled matrix of 27 unique CVs to isolate the impact of document length and content density on model vulnerability.

Dataset Composition

The dataset is categorized by two primary dimensions: **Professional Seniority** (correlating with document length) and **Achievement Density (AD)**. For each intersection of these dimensions, three distinct CVs were authored to ensure results are not artifact-specific, totaling 27 documents.

A. Seniority

The Seniority dimension is related with the length of the document, i.e. in general an entry level profile has a shorter CV than a long term professional. Document's length serves as a proxy for evaluating

g the LLM's attention span and context window management.

- **Junior (J):** 1-page documents (~400–600 words) representing entry-level profiles with high information visibility.
- **Mid-Level (M):** 2–3 page documents (~1,000–1,500 words) reflecting standard professional complexity.
- **Senior (S):** 4–6 page documents (~2,500+ words) providing extensive "noise" where malicious prompts can be deeply embedded.

B. Achievement Density (AD) Classes

An **achievement** is defined as a discrete, quantifiable result—typically containing a metric (percentage, currency, or time-frame, "*manage 1M\$ budget*")—or a specific, high-impact milestone (e.g., "*Led a team of 50*" or "*Received the Innovation Award*").

We define **Achievement Density (AD)** of a CV as the frequency of quantifiable, metric-driven accomplishments declared in the CV relative to the total word count, W_{total} .

$$AD = \frac{N_{achievements}}{W_{total}} \times 100$$

where $N_{achievements}$ represents the count of distinct achievements, and W_{total} is the total word count of the document.

This metric allows us to evaluate whether high concentrations of "legitimate signal" help obfuscate "malicious instructions."

To ensure a balanced experimental distribution, we established the following **Density Classes**:

- **High AD:** $AD \geq 1.5$ (highly condensed, data-driven profiles, high frequency of KPIs and quantifiable results);
- **Medium AD:** $0.5 < AD < 1.5$ (Standard professional resumes with a balanced mix of responsibilities and achievements);
- **Low AD:** $AD \leq 0.5$ (narrative-heavy or entry-level profiles where achievements are sparse relative to descriptive text).

The following table shows the dataset composition of the actual experimented 21 CVs along the two dimensions Seniority-Documents Length and Achievement Density.

Category	AD=Achievement Density	CV IDs	Target Role	Avg. Word Count	Key CV's Characteristics
Junior	High	CV01, 02, 03	Jr. Developer	450	Metric-heavy intern/grad projects.
Junior	Med	CV04, 05, 06	Marketing Asst	550	Balanced task/achievement list.
Junior	Low	CV07, 08, 09	Office Admin	600	Descriptive, narrative-style duties.
Mid-Level	High	CV10, 11, 12	Project Manager	950	Heavy on budget/KPI percentages.
Mid-Level	Med	CV13, 14, 15	Sales Rep	1100	Mix of quotas and client relations.
Mid-Level	Low	CV16, 17, 18	HR Specialist	1200	Long descriptions of processes.
Senior	High	CV19, 20, 21	CTO / Director	2200	Dense list of patents/M&A/scaling.
Senior	Med	CV22, 23, 24	Consultant	2500	Case studies mixed with bio.
Senior	Low	CV25, 26, 27	Academic Prof.	3000	Verbose history, few numeric KPIs.

Injection Protocol

Each of the 27 documents was subjected to the four injection prompt types:

- *p.Pos*
- *p.Neg*
- *p.Pos-Neg*
- *p.Pos-Med*

each prompt is contained in its respective text file in the **PROMPT_INJECTION** folder for a total experimental corpus of 108 attack scenarios to submit to each LLM.

The CV instance with of an injection prompt type may require adaptation to some specific CV content (e.g. Firstname/Lastname of the CV owner).

To account for stochastic variations in model output, each scenario should be executed for at least 10 independent trials ($N=10$) with the temperature parameter set to 0.

How to generate the Plain and Injected CVs Dataset

Generate the plain CVs dataset:

- 1) For each CV contained in the **CV** folder
 - 2) remove from its content each label in square bracket [Achievement X] where X represent the achievement id number in the CV
- (See example for CV01 in the **SAMPLE_PLAIN_INJECTED_CV01** folder)

Generate the single CV injection type *p.type*:

- 1) given a CVXX in the plain dataset where XX is the CV identifier integer
 - 2) search in the **PROMPT_INJECTION** folder for the *p.type* text file
 - 3) generate personalized version for CVXX , e.g. mentioning Firstname/Lastname
 - 4) hide the prompt in CVXX content, e.g. insert in header as text with color white on white background.
- (See examples for CV01 in the **SAMPLE_PLAIN_INJECTED_CV01** folder)

USER PROMPTS

The following user prompts text files are provided in the **USER_PROMPTS** folder to test regular single CV (**u1**) or multiple CVs evaluations (**u2**), after loading the corresponding CV, either in a separate LLM chat, for Single CV, or in the same LLM chat for multiple CVs. The "prompt extended with prevention measures" versions, (**u1.X** and **u2.X**) are meant to evaluate the efficacy of countermeasure to prevent prompt injection;

- u1_Regular_SINGLE_CV_Prompt.txt
- u2_Regular_MULTIPLE_CVs_Prompt.txt

- u1_X_SINGLE_CV_Prompt_with_Prevention_Measures.txt
 - u2_X_MULTIPLE_CV_Prompt_with_Prevention_Measures Measures.txt
-

DATASET Terms of Use

License Grant

This dataset is made available under a [Creative Commons Attribution 4.0 International License \(CC BY 4.0\)](#)."

Permitted Use

Users are permitted to copy, distribute, display, and create derivative works for both non-commercial and commercial purposes, provided that appropriate attributions is given as described in the next paragraph.

The dataset does not contain sensitive restricted or personal data of any kind.

Required Attribution

When results of products using the dataset are described in publications and/or report of any nature (e.g. scientific, technical, business etc.), when derived datasets are generated from, and/or software application are developed to manage the dataset or its derivatives then appropriate credit must be given, in the form of bibliographic reference, and/or acknowledgement, and/or in the description or product, software or derived dataset and source code comments by:

- *providing a link to the license,*
- *indicating if changes were made*
- *citing the dataset and related journal article as:*

Alfredo Milani, Valentina Franzoni, and Emanuele Florindi. 2026. "Indirect Prompt Injection CV Dataset." Harvard Dataverse. <https://doi.org/10.7910/DVN/8YBJQ4>.

Alfredo Milani, Valentina Franzoni, Emanuele Florindi, Indirect Prompt Injection in Large Language Models, (2026) , Neural Computing and Application, Springer

Limitation of Liability & Warranties

The dataset is provided 'as it is' without warranty of any kind. The creators assume no liability for errors, omissions, or any damages arising from the use of this data.
